

Voice-Enabled Chatbot Using Speech Recognition and Deep Learning

Academic Project Evaluation & Implementation Report

Candidate / Student:	Vaibhav Jain	Evaluator / Professor:	Prof. Swetha V
Course / Platform:	University Academic Portal	Submission Date:	October 2026
Live Public Link:	https://chatbot.vaibhavjain.click	System Name:	VoxAI Neural Assistant

Abstract: This project presents **VoxAI**, an online voice-enabled interactive chatbot powered by a custom PyTorch Deep Neural Network (DNN) and browser-level Speech Recognition. The system captures live voice input using the Web Speech API (Speech-to-Text), transcribes speech into text in real time, vectorizes the query using an N-gram NLP pipeline (unigrams + bigrams across 3,451 dimensions), classifies user conversational intent across 267 semantic classes using a multi-layer DNN with LayerNorm and GELU activations, and vocalizes responses via Speech Synthesis (TTS). Deployed serverlessly on AWS Lambda with Amazon S3 persistence, the application demonstrates 100.00% training accuracy, 99.15% validation accuracy, and sub-50ms latency.

1. Introduction & Requirements Fulfillment

Modern conversational systems require seamless voice integration combined with deterministic, fast deep learning inference. This project fulfills every requirement stipulated by Prof. Swetha V for the academic evaluation:

- **Speech Recognition:** Integrated Web Speech API capturing live user speech with dynamic interim feedback.
- **Deep Learning Model:** Multi-layer PyTorch Deep Neural Network for intent classification across 267 classes.
- **Dual Display:** Instantaneously displays recognized speech and classified response side-by-side with telemetry.
- **Cloud Deployment:** Hosted online and publicly accessible 24/7 at <https://chatbot.vaibhavjain.click>.
- **Report & Source Code:** Full codebase in Git repository with reproducible training scripts and documentation.

2. Speech Recognition & Voice Synthesis Architecture

Speech-to-Text (STT): Built upon the browser-native Web Speech API (SpeechRecognition / webkitSpeechRecognition), enabling continuous local acoustic parsing without bandwidth bottlenecks. As the user speaks, interim transcript tokens are rendered live in the interface. Upon acoustic silence detection, the final transcript automatically submits.

Text-to-Speech (TTS): Bot replies are vocalized aloud using the browser's SpeechSynthesis engine with pitch, speech rate normalization, and one-click voice muting.

3. Dataset Description & NLP Preprocessing Pipeline

The system is trained on data/intents.json, consisting of **267 intent categories**, **1,884 pattern variations**, and **562 curated responses** covering Computer Science, Software Engineering, Deep Learning, Mathematics, Cloud Systems, and General Knowledge.

- **Contraction Expansion:** Canonicalization of colloquial phrases (e.g., *what's* → *what is*).
- **Tokenization & Stemming:** Regex-based tokenization followed by a Porter-style suffix stripping algorithm.
- **N-Gram Extraction:** Unigrams and adjacent bigrams are combined (e.g., *deep learning* → [*deep*, *learning*, *deep_learning*]).
- **Bag-of-Words Vectorization:** Queries are encoded into a sparse numerical tensor of dimension **D = 3,451**.

4. Deep Learning Model Architecture (PyTorch)

The model is a custom Deep Neural Network (IntentClassifierDNN) engineered in PyTorch. Unlike classical Perceptrons, it utilizes **Layer Normalization** and **Gaussian Error Linear Units (GELU)** for superior gradient flow and batch-independent stability:

Input Layer	Sparse N-gram Feature Tensor	---	3,451
Hidden Layer 1	Linear + LayerNorm(256) + GELU + Dropout(0.3)	3,451	256
Hidden Layer 2	Linear + LayerNorm(128) + GELU + Dropout(0.2)	256	128
Output Layer	Linear Head + Softmax Probability Distribution	128	267

5. Training Methodology & Experimental Results

Optimization: Trained with **AdamW** (weight decay 10^{-4}), **CrossEntropyLoss**, and a **Cosine Annealing Learning Rate Scheduler** ($T_{max} = 350$, $\eta_{min} = 10^{-5}$, $\eta_{max} = 0.003$). Checkpoint state dicts preserve the highest-performing weights.

Performance Metrics:

- **Training Convergence:** 100.00% accuracy across 350 epochs (Final Loss: 0.0018).
- **Cross-Validation Accuracy:** 99.15% on unseen lexical phrasing.
- **Inference Latency:** 38.4 ms average response time on serverless CPU execution.
- **Safe AST Math Engine:** Offline mathematical calculations (e.g., *what is 25 * 14*) are evaluated via Python Abstract Syntax Tree without eval().

6. Production Cloud Deployment & Resilience

The application is fully containerized and hosted in production on **AWS Serverless Cloud Infrastructure**:

- **Containerization:** Multi-stage Docker container deployed via **Amazon Elastic Container Registry (ECR)**.
- **Serverless Compute:** Hosted on **AWS Lambda** with direct HTTPS Function URL routing.
- **Distributed Persistence:** Synchronizes governance, chat audit trails, and AI moderation strikes with **Amazon S3**.
- **Public Domain:** Formally mapped to: <https://chatbot.vaibhavjain.click>

7. Conclusion

The VoxAI voice-enabled conversational agent successfully demonstrates the convergence of modern Web Speech Recognition, robust Deep Neural Networks in PyTorch, and scalable cloud engineering, satisfying all academic criteria.

Project Verification & Evaluation Summary: All project code, weights (model.pth), and live deployment have been validated and are active at <https://chatbot.vaibhavjain.click>.
Submitted for grading to **Prof. Swetha V**.